STATEMENT OF WORK · CONFIDENTIAL

AI-Powered Call Analytics & Compliance Intelligence Platform

For Call Monitoring & Agent Performance Management

DOCUMENT TYPE	PREPARED FOR	VERSION	DATE	STATUS
Statement of Work	ICICI Bank	2.1	May 2026	Draft — Awaiting Review

EXECUTIVE SUMMARY

The Problem

Banks operating EMI recovery and loan collection teams deploy large numbers of tele-callers who interact with customers daily. These calls are recorded — yet the recordings remain largely unanalysed at scale, creating a critical blind spot across three dimensions.

Unverified Dispositions

Tele-callers self-report call outcomes. There is currently no systematic way to verify whether those reports accurately reflect what was actually said during the call.

No Script Compliance Visibility

The bank provides a prescribed calling script, but there is no mechanism to measure how closely agents follow it — or to identify which stages are consistently skipped or handled poorly.

Missed Customer Intelligence

Customers often reveal the real reasons they cannot pay — financial hardship, disputed amounts, employment changes — but this information is lost if not captured and analysed systematically.

> This platform resolves all three gaps by ingesting call recordings and tele-caller dispositions, transcribing every call with full speaker separation, and applying structured AI analysis to produce per-call intelligence reports — at production scale, overnight.

SECTION 2

Objectives

Five measurable outcomes the system is designed to deliver for the bank's operations and compliance teams.

Disposition Verification

Cross-check tele-caller submitted dispositions (D0) against AI-inferred dispositions (D1) derived from actual call content. Flag every mismatch with severity and explanation.

Agent Quality Scoring

Score each call against the bank's prescribed script, stage by stage, on three dimensions: flow correctness, response handling, and call quality.

Customer Intelligence

Extract structured reasons for non-payment or deferral — financial hardship, disputes, medical emergencies — from each call's transcript.

Operational Feedback

Provide per-call and aggregate coaching feedback highlighting what agents did well, where they failed, and how specific moments could have been handled better.

Audit Trail

Maintain a complete, traceable record of every call processed, every disposition mismatch, and every quality score — queryable by call, agent, date, or outcome.

SECTION 3

Scope of Work

✓ IN SCOPE	— OUT OF SCOPE
✓ Secure data ingestion from bank-provided SFTP servers or cloud storage buckets	— Real-time call interception or live-call monitoring
✓ Audio file preprocessing and format normalisation	— CRM integration or loan management system changes on the bank's side
✓ ASR with full speaker diarization (Agent vs. Customer separation)	— Legal or compliance advisory on banking regulations
✓ LLM-based structured analysis via vendor Batch API	— Human review or auditing of flagged calls (bank-side responsibility)
✓ Generation of per-call structured JSON reports	— Training or fine-tuning of any AI model on bank data
✓ Disposition comparison (D0 vs D1) and mismatch flagging	
✓ Agent performance scoring and coaching feedback generation	
✓ Structured output delivery back to the bank (API, SFTP push, or dashboard)	
✓ Scalable pipeline designed for 100GB / 2,50,000 calls per day	

SECTION 4

Input Specification

The bank provides two categories of data per call. Both must be present and matched by `call_id` before processing begins.

4.1 Audio Files (Call Recordings)

ATTRIBUTE	DETAIL
Format	`.mp3`, `.wav`, `.ogg`, `.m4a` — all common telephony formats supported
Duration	Typically 1–15 minutes per file
Channels	Mono or stereo (stereo preferred for cleaner diarization)
Language	Hindi, English, or Hinglish (mixed) — all handled natively
Naming Convention	`{call_id}_{agent_id}_{date}_{time}.ext` — to be finalised during onboarding
Encoding	8 kHz–16 kHz sample rate (telephony standard); higher rates accepted. Audio is normalised to 16 kHz mono before STT submission for consistent vendor output.

4.2 Call Disposition (Tele-caller Submitted)

Structured metadata submitted by the tele-caller after each call — their own characterisation of the call outcome. This becomes D0 in the disposition verification module.

FIELD	TYPE	DESCRIPTION
`call_id`	String	Unique identifier — links this disposition record to its audio file
`agent_id`	String	Tele-caller employee ID
`customer_id`	String	Loan account or borrower identifier
`disposition_code`	Enum	PAID PROMISE_TO_PAY REFUSED NOT_REACHABLE DISPUTED
`disposition_notes`	Text	Free-text notes from the agent about the call outcome
`call_datetime`	Timestamp	ISO 8601 format
`loan_type`	String	EMI / Personal Loan / Credit Card / etc.
`overdue_amount`	Float	Amount overdue (in INR) at time of call

> Delivery format for dispositions — CSV, JSON, or direct API push — to be agreed during onboarding. Both audio and disposition must share the same `call_id` for the system to pair and process them.

SECTION 5

Data Ingestion & Transfer

The bank may deliver data via either of two secure transfer methods. The system supports both concurrently.

Option A — SFTP

Secure File Transfer Protocol

- Bank drops audio files and disposition CSV/JSON to a designated SFTP path
- Ingestion service polls at a configurable interval (e.g., every 15 minutes or nightly)
- SSH key-based authentication only — password auth disabled
- Processed files acknowledged and archived on the SFTP after receipt

Option B — Cloud Storage

S3 / Azure Blob / GCS

- Bank uploads to an agreed cloud storage bucket in their preferred provider
- Object-creation events trigger the ingestion service in near real-time
- IAM-scoped, least-privilege credentials exchanged during onboarding
- No persistent connection required — event-driven, zero polling overhead

Ingestion Processing Steps

STEP 1
Format Validation — Verify audio codec, duration range, and disposition schema

STEP 2
Deduplication — Each `call_id` checked against a registry; duplicates silently dropped with a log entry

STEP 3
call_id Pairing — Audio matched to its disposition record; unmatched held for up to 24 hours

STEP 4
Validation Complete — Call confirmed paired, formatted, and ready for processing

STEP 5
Queue Enqueue — Validated, paired calls placed in the message queue for STT dispatch

Arrival Guarantees

EXACTLY-ONCE

Each `call_id` is tracked in a deduplication registry. Duplicate arrivals are detected and discarded with a log entry — no double-processing.

FAILED INGESTION

Files failing integrity or format checks are quarantined with a reason code. A daily error report is generated for the bank's review.

OUT-OF-ORDER DELIVERY

Audio and disposition may arrive independently. The system holds the earlier-arriving piece for up to 24 hours before flagging the pair as incomplete.

SECTION 6

Processing Pipeline

Each call passes through two sequential processing stages. Both are powered by external vendor APIs — no on-premise GPU infrastructure is required.

1 Stage 1 — Speech Processing: ASR + Diarization

Objective: Convert raw audio into a clean, speaker-labelled JSON transcript ready for LLM analysis.

AUDIO PRE-PROCESSING
Noise reduction, volume normalisation, format conversion to 16 kHz mono WAV for consistent vendor input

SPEAKER DIARIZATION
Separates agent and customer speaker turns. Assigns labels: `[AGENT]` / `[CUSTOMER]`

ASR (SPEECH-TO-TEXT)
Hindi / English / Hinglish transcription via vendor API. Produces timestamped utterances per speaker turn

OUTPUT ARTIFACT
Structured JSON transcript stored in the Transcript Store and passed to Stage 2

Output format: `[{"speaker":"AGENT","start":0.0,"end":4.2,"text":"..."},{"speaker":"CUSTOMER","start":4.5,"end":9.1,"text":"..."},...]`

2 Stage 2 — LLM Analysis

Objective: Apply structured intelligence to each transcript using a large language model, producing four analysis modules per call.

LLM INPUTS
Full diarized transcript + bank's calling script (versioned) + tele-caller's D0 disposition

PROMPT ENGINEERING
System prompt encodes the calling script and instructs the model to produce structured JSON output for each module

STRUCTURED OUTPUT
Model returns deterministic JSON for all 4 modules: Disposition, Refusal, Scoring, Intelligence

REPORT STORAGE
Per-call JSON report written to the Reports Database, indexed by `call_id`, `agent_id`, and date

> The calling script is provided by the bank once, versioned, and referenced in the LLM system prompt. Script updates require advance notice to allow prompt recalibration.

SECTION 7

Data Flow Architecture

The complete end-to-end flow from bank data delivery to report availability. The system is API-native and queue-driven — no GPU infrastructure required on our side.

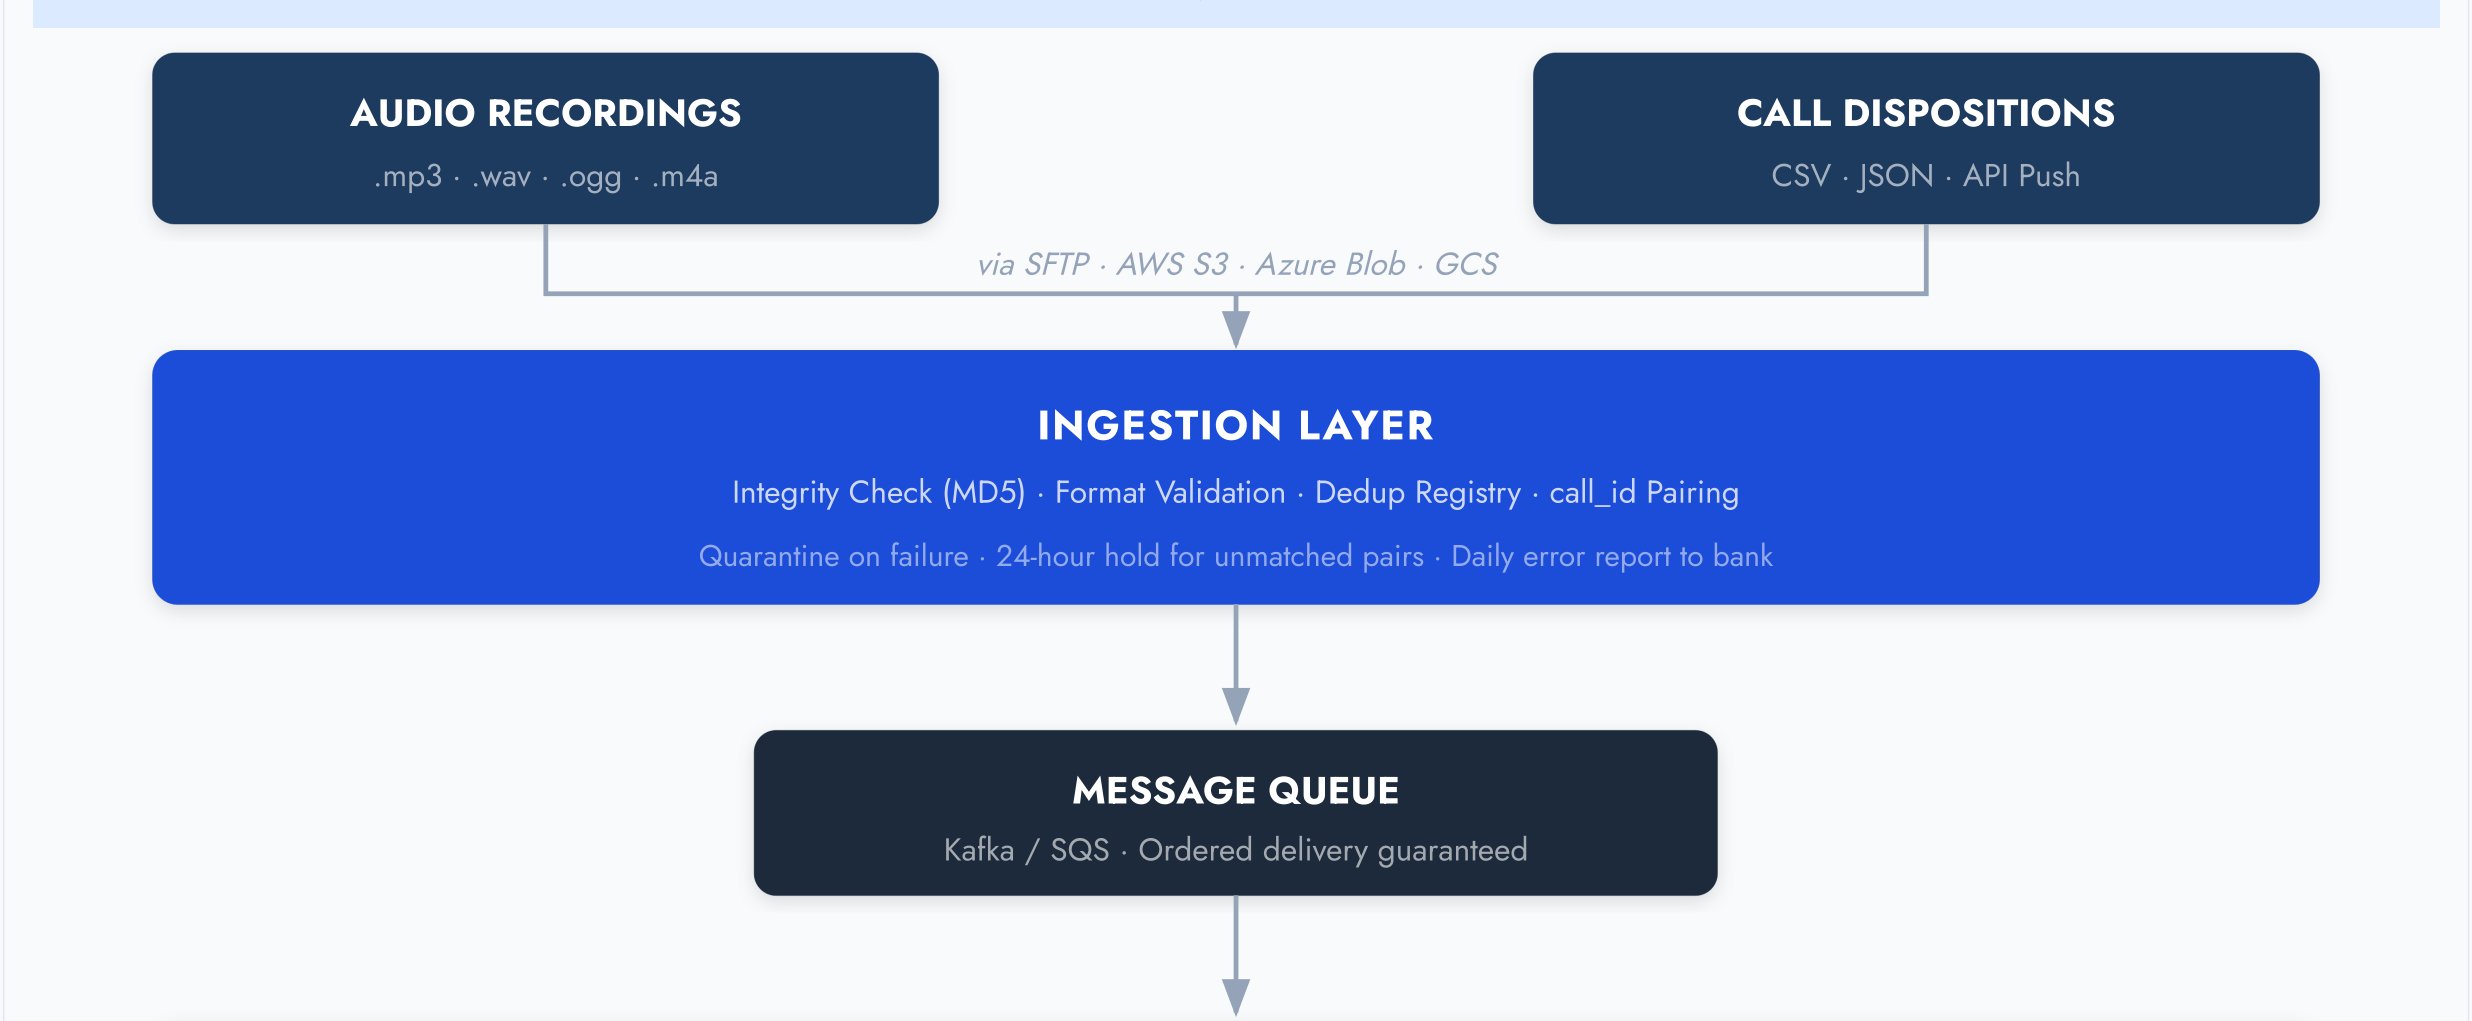

BANK SIDE

AUDIO RECORDINGS
.mp3 · .wav · .ogg · .m4a

CALL DISPOSITIONS
CSV · JSON · API Push

via SFTP · AWS S3 · Azure Blob · GCS

INGESTION LAYER
Integrity Check (MD5) · Format Validation · Dedup Registry · call_id Pairing
Quarantine on failure · 24-hour hold for unmatched pairs · Daily error report to bank

MESSAGE QUEUE
Kafka / SQS · Ordered delivery guaranteed

STT DISPATCH LAYER — VENDOR API
100 Concurrent Jobs · Async Submit-and-Poll · Audio Preprocessing · Speaker Diarization
Auto-failover to secondary vendor · Retry up to 3× on job failure · Dead-letter queue on sustained errors

TRANSCRIPT STORE
Diarized JSON · Speaker-labelled · Timestamped Utterances

LLM ANALYSIS — BATCH API
Nightly Batch Run · 4 Analysis Modules per Call · Structured JSON Output
~3.5 hours to process 2,50,000 calls

REPORTS DATABASE
Per-call JSON · Queryable · Audit Trail

REST API
Query by call_id / date / agent

SFTP / CLOUD PUSH
JSON / CSV to bank destination

QA DASHBOARD
Browse, filter & export

SECTION 8

LLM Analysis — Output Specification

Each processed call produces a structured report composed of four analysis modules. The LLM receives the diarized transcript, the bank's calling script, and the tele-caller's D0 disposition — and returns deterministic JSON for each module.

1 Disposition Verification

The LLM independently infers the true call outcome (D1) from transcript content and compares it against the tele-caller's submitted disposition (D0). Every mismatch is recorded with severity and explanation.

OUTPUT FIELD	TYPE	DESCRIPTION
`ai_disposition` (D1)	Enum	PAID PROMISE_TO_PAY REFUSED NOT_REACHABLE DISPUTED PARTIAL_COMMITMENT UNCLEAR
`bank_disposition` (D0)	Enum	Original submission from the tele-caller — unchanged
`disposition_match`	Boolean	Whether D0 and D1 agree
`mismatch_severity`	Enum	NONE MINOR SIGNIFICANT CRITICAL
`mismatch_explanation`	Text	When mismatch exists: what the transcript shows vs. what was reported
`confidence_score`	Float 0–1	LLM confidence in its inferred disposition

> **Example mismatch:** Transcript shows customer said "dekh lenge" (we'll see) with no specific date commitment. LLM infers D1 = `UNCLEAR`. Tele-caller submits D0 = `PROMISE_TO_PAY`. Mismatch severity = `SIGNIFICANT`. The customer isn't committing, they're deflecting. The agent logged PROMISE_TO_PAY, but the transcript contains no amount, no date, and no actual agreement — so the LLM infers UNCLEAR. The mismatch is SIGNIFICANT because the bank may now deprioritise this account and wait for a payment that won't arrive.

2 Reason for Refusal / Non-Payment

When the customer is unwilling or unable to pay, the LLM extracts the primary and secondary reasons as expressed in the conversation, along with sentiment and escalation risk.

OUTPUT FIELD	TYPE	DESCRIPTION
`payment_status`	Enum	AGREED PARTIAL DEFERRED REFUSED UNCLEAR
`primary_reason`	Enum	FINANCIAL_HARDSHIP · DISPUTES_AMOUNT · ALREADY_PAID · WILL_PAY_LATER · JOB_LOSS · MEDICAL_EMERGENCY · UNRESPONSIVE · OTHER
`reason_verbatim`	Text	Direct quote or close paraphrase from the transcript supporting the reason
`secondary_reasons`	Array	Additional reasons mentioned, if any
`customer_sentiment`	Enum	COOPERATIVE NEUTRAL AGITATED HOSTILE
`escalation_risk`	Enum	LOW MEDIUM HIGH — likelihood of dispute or regulatory escalation

3 Script Adherence & Agent Quality Score

The LLM evaluates the agent's performance against the bank's prescribed calling script — stage by stage — and generates a three-dimensional quality score with actionable coaching feedback.

3A — AGGREGATE SCORE DIMENSIONS

DIMENSION	SCALE	WHAT IT MEASURES
`flow_correctness`	0–10	How faithfully the agent followed the correct sequence of script stages
`response_handling`	0–10	How well the agent handled objections, refusals, and customer emotions
`call_quality`	0–10	Overall communication clarity, tone, and professionalism
`composite_score`	**0–10**	**Weighted average of all three — primary KPI for agent performance tracking**

3B — STAGE-BY-STAGE ASSESSMENT FIELDS

FIELD	TYPE	DESCRIPTION
`stage_id`	String	e.g., A, B, C, D, E — maps to stages in the bank's calling script
`stage_name`	String	e.g., Identity Verification, Overdue Amount Statement, Consequence Nudge
`status`	Enum	COMPLETED PARTIALLY_COMPLETED SKIPPED IMPROVISED
`observation`	Text	What the agent actually did at this stage
`deviation_note`	Text	If not COMPLETED: what was missing or incorrect

3C — COACHING FEEDBACK FIELDS

FIELD	TYPE	DESCRIPTION
`strengths`	Array[Text]	2–3 specific things the agent handled well
`improvement_areas`	Array[Text]	2–3 specific areas where the agent fell short
`suggested_alternative`	Text	A brief example of how a missed or poor moment could have been handled better

ILLUSTRATIVE SAMPLE OUTPUT

FLOW CORRECTNESS	RESPONSE HANDLING	CALL QUALITY
6/10	7/10	8/10

STAGE-BY-STAGE ASSESSMENT

Stage A ✓ Identity confirmed correctly. Good use of repair phrase before re-asking the verification question.

Stage B ✓ Loan number and overdue amount stated accurately. Same-day payment ask delivered cleanly.

Stage C ✓ Consequence nudge (penalty + credit health) correctly applied after the customer's first refusal.

Stage E ✗ Customer said "loan settlement karo" — a firm refusal. Script mandates Stage E (delay reason) before closing. Agent skipped Stage E entirely and went straight to SMS link. Missing: *"What is the specific reason you cannot pay today?"*

Settlement ~ Redirecting to branch/customer care is an acceptable improvisation, but caused an early exit from the script flow.

4 Overall Call Intelligence

A holistic synthesis of the call — not limited to script compliance. Captures the broader picture of what happened, what it means for collections strategy, and what the bank should do next.

OUTPUT FIELD	TYPE	DESCRIPTION
`call_summary`	Text	3–5 sentence neutral summary of the entire call
`key_pain_points`	Array[Text]	Specific pain points raised by the customer that may affect payment likelihood
`agent_missed_opportunities`	Array[Text]	Moments where a different response could have meaningfully changed the outcome
`notable_observations`	Array[Text]	Anything unusual, potentially non-compliant, or noteworthy in the call
`recommended_next_action`	Text	Suggested next step for the bank's collections team on this account
`overall_feedback`	Text	2–3 sentence integrated assessment of collections effectiveness

SECTION 9

System Output & Deliverables

For every processed call, the system produces a unified JSON report — and delivers it back to the bank through the agreed channel.

Per-Call Report Structure

```
{
  "call_id":              "CALL_20260505_0981",
  "agent_id":             "AG_447",
  "processed_at":         "2026-05-05T14:32:00Z",
  "transcript":           { speaker-labelled turns... },
  "disposition_verification": { D0, D1, match, severity, explanation },
  "refusal_analysis":       { payment_status, reason, sentiment, risk },
  "agent_quality":          { scores, stage_assessment, coaching_feedback },
  "call_intelligence":      { summary, pain_points, next_action, feedback }
}
```

Delivery Channels

CHANNEL	DESCRIPTION
SFTP / Cloud Push	JSON or CSV reports delivered to the bank's designated SFTP path or cloud bucket after nightly processing
QA Dashboard	Web-based interface for QA managers to browse, filter, score-sort, and export call reports
Aggregate Reports	Daily / weekly summaries: mismatch rates, average scores by agent and team, top refusal reason categories

SECTION 10

Infrastructure & Scale

The platform is built entirely around external vendor APIs for STT and LLM. No GPU infrastructure is required. The architectural challenge is intelligent API orchestration — maximising throughput within rate limits at the lowest operational overhead.

10.1 Volume at 100GB / Day

100GB	~2.5L	~0.4 MB	~25 sec	~1 GB	~0.2 GB
AUDIO PER DAY	CALLS PER DAY	AVG FILE SIZE	AVG CALL DURATION	TRANSCRIPTS / DAY	REPORTS / DAY

> Key principle: once audio is transcribed, it is never read again by the system. All downstream LLM processing, re-analysis, and report queries operate on the stored transcript JSON (~4 KB), not the 400 KB audio file. Re-running analysis with an updated prompt costs almost nothing.

10.2 STT — Parallel Concurrent Dispatch

STT vendor APIs operate on an **async submit-and-poll model** — you POST a file URL, receive a job ID, and poll for completion. This is inherently parallelisable. The STT Dispatch Layer maintains a pool of **100 concurrent in-flight jobs** at all times. As each job completes, the slot is freed and the next call is immediately submitted — the pool stays saturated, the system never waits.

BATCH API VS SYNC REST

The 25-second average implies a distribution — a significant portion of files will exceed 30 seconds and take the async batch path (with polling overhead), not the sync REST path. Send everything through the batch API regardless of duration since the average is 25 sec, the distribution likely skews short, making a hybrid approach unnecessarily complex. Batch API is simpler to manage and never slower than sync

PER-JOB TIMING — WHAT HAPPENS INSIDE EACH API CALL

For a 25-second audio file at 0.4 MB, each vendor job goes through three steps in sequence. These are short calls — vendor compute is minimal, so API overhead is now the dominant cost:

STEP 1 · FILE UPLOAD
~0.1 sec
0.4 MB over a high-throughput cloud connection — negligible transfer time at scale

STEP 2 · STT COMPUTE
~0.25 sec
Vendor processes at ~100× real-time. 25 sec audio ÷ 100 = ~0.25 sec of actual compute

STEP 3 · API OVERHEAD
~2.0 sec
HTTP round trips, vendor queue wait, poll cycle latency — irreducible floor

THROUGHPUT CALCULATION

STT — DERIVATION

Avg wall-clock time per job	`0.1 + 0.25 + 2.0 sec`	`= 2.35 sec / job`
Effective throughput at 100 concurrent	`100 slots ÷ 2.35 sec`	`= 42.5 completions / sec`
Total time for 2,50,000 calls	`2,50,000 ÷ 42.5`	`= 5,882 sec = 98 min`
STT Estimated Time — 2,50,000 calls at 100 concurrent jobs		`≈ 1.63 hours`

CONCURRENCY IMPACT — STT STAGE (2,50,000 CALLS)

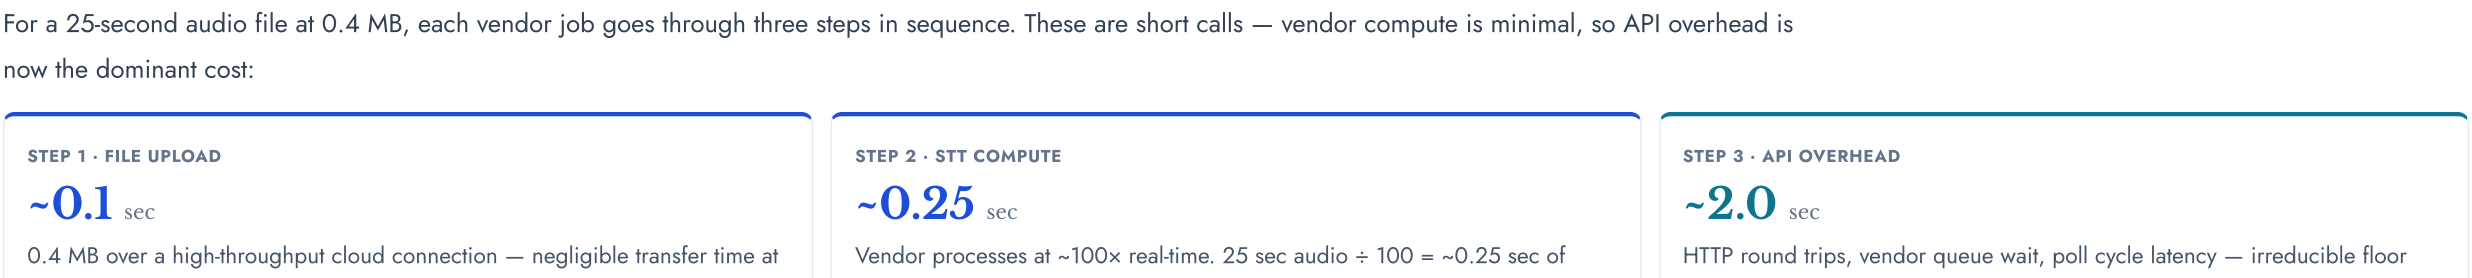

> At 100 concurrent STT jobs, all 2,50,000 calls are fully transcribed comfortably under **2 hours**. These short 25-second calls process fast because vendor compute is minimal — high concurrency is the critical lever, not per-call processing speed.

10.3 LLM — Batch API Processing

Both Anthropic and OpenAI offer a dedicated **Batch API** designed for high-volume, non-time-critical workloads. The Batch Assembler collects all transcripts once STT completes and submits them as a single batch — the vendor processes asynchronously on their infrastructure and returns a complete result file with no persistent workers required.

TOKEN BREAKDOWN PER CALL

Each call submits one self-contained JSON object to the Batch API. These are short calls — the system prompt (calling script + output schema) is the dominant cost component per request:

SYSTEM PROMPT
~500 tokens
Calling script + 4-module JSON schema + instructions. Constant across all calls.

TRANSCRIPT INPUT
~100 tokens
25 sec speech ≈ 62 words ≈ ~100 tokens of diarized, speaker-labelled text

D0 + METADATA
~50 tokens
Disposition code, agent notes, loan type, overdue amount

OUTPUT — 4 MODULES
~550 tokens
4 modules × ~140 tokens each — all verdicts, scores, and coaching feedback

TOTAL PER CALL
~1,200 tokens
650 tokens input + 550 tokens output

BATCH PROCESSING CALCULATION

LLM BATCH API — DERIVATION

Tokens per call	`650 input + 550 output`	`= 1,200 tokens`
Total tokens (2,50,000 calls)	`2,50,000 × 1,200`	`= 300M tokens`
Batch API throughput (vendor midpoint)	`75M – 100M tokens / hour`	`≈ 87.5M / hr`
Processing time	`300M ÷ 87.5M`	`= 3.43 hours`
LLM Estimated Time — 2,50,000 calls via Batch API		`≈ 3.5 hours`

BATCH API — NIGHTLY PROCESSING AT A GLANCE

Total tokens (2,50,000 calls)	**~300 million tokens**
Batch API throughput	**~87.5M tokens / hour**
LLM processing time	**~3.5 hours**
Cost vs. individual sync calls	**Significantly reduced**
Infrastructure required	**Single lightweight nightly job**
Result delivery	**Single downloadable file — all 2,50,000 reports**

Each call in the batch is submitted as a self-contained JSON object with a `custom_id` matching the `call_id`, ensuring unambiguous result mapping. STT and LLM run **sequentially** — the batch assembler waits until all transcripts are ready, then submits the full set at once.

END-TO-END PIPELINE SUMMARY

2,50,000 CALLS / DAY · SEQUENTIAL STAGES

STT Transcription 100 concurrent async jobs · 42.5 completions / sec · 2,50,000 ÷ 42.5 = 5,882 sec	~1.63 hrs
LLM Analysis Nightly Batch API · 300M tokens · 300M ÷ 87.5M = 3.43 hrs · runs after STT completes	~3.5 hrs
Total End-to-End Last audio file in → all 2,50,000 reports available	~5.1 hrs

> All 2,50,000 calls — transcribed, scored, and fully analysed — within **~5.1 hours**. Well within a single overnight window. The bank's QA team receives a complete, annotated report set each morning with zero manual intervention.

10.4 Data Retention Policy

Audio files are large (100GB/day) and have no operational value after transcription. Transcripts are small (~40 KB each) and carry all analytical value.

Audio Files

7-day retention

Raw audio files are automatically deleted after 7 days. Once transcribed, the audio file is never accessed again — all re-analysis uses the stored transcript.

Transcripts & Reports

Long-term storage

Diarized transcripts and structured analysis reports are retained indefinitely (or per the bank's compliance requirements). These are the primary analytical assets of the system.

10.5 Fault Tolerance & Reliability

STT job failure	Automatic retry up to 3 times per job. On sustained failure, moved to dead-letter queue and flagged in daily error report.
STT vendor outage	Automatic failover to a pre-configured secondary STT vendor. The dispatcher switches transparently — no manual intervention.
LLM batch partial failure	Batch API returns per-item success/failure. Failed items are logged, flagged in the daily error report, and automatically requeued for the next nightly batch run.
LLM batch delayed	If batch results are not received within the expected window, the ops team is alerted immediately and the vendor is contacted for status. No data is lost — all transcripts remain in store and will be reprocessed.
Duplicate file delivery	`call_id` deduplication registry prevents double-processing. Duplicates are silently discarded with a log entry.
Queue worker crash	Jobs are acknowledged only after a successful write to the transcript store. Unacknowledged jobs are automatically requeued — no data loss.
SFTP / cloud downtime	Retry with exponential backoff. Alert triggered after 3 consecutive failures. No data is lost — the bank's source remains intact.

SECTION 11

Security & Data Compliance

Call recordings contain sensitive customer data. The following controls are applied across all layers of the system.

Data in Transit

- › All transfers (SFTP, cloud sync, API calls) use TLS 1.2 or higher
- › SFTP connections use SSH key-based authentication or password auth
- › Cloud storage access governed by IAM least-privilege policies

Data Handling & Locality

- › Customer PII is not extracted or stored separately — remains embedded in transcripts only
- › Audio files stored under `call_id` reference only — no additional customer identity mapping
- › Data stored within Indian data centres unless otherwise agreed

LLM Data Privacy

- › Transcripts processed via private API endpoints — not shared consumer models
- › No training of any model on bank data without explicit written consent
- › On-premise or VPC-deployed LLM available for banks with strict data locality requirements

SECTION 12

Assumptions & Dependencies

The following are assumed to be true for the system to operate as described. Deviations from these assumptions will require design review.

1. The bank will provide the calling script in a structured, versioned format (PDF or plain text). Script changes will be communicated at least two weeks in advance, with a version number, to allow prompt recalibration.
2. Audio files will include both agent and customer audio — not agent-only recordings. Diarization requires both speaker channels to be present.
3. Each audio file maps 1:1 with a call disposition record via a shared `call_id`. The naming convention and `call_id` format will be agreed during onboarding.
4. The bank's SFTP server or cloud storage bucket will be available and accessible with agreed credentials throughout operations.
5. Call recordings are in a standard telephony format (MP3, WAV, OGG). Non-standard formats may require additional preprocessing lead time.
6. The bank will designate a technical point of contact for onboarding, credential exchange, and ongoing issue resolution.
7. Disposition codes and categories are agreed upfront and remain stable. New categories will be introduced with at least two weeks' notice.
8. The bank accepts that AI-inferred disposition (D1) is probabilistic in nature — it is a quality signal and audit tool, not a legal record or primary disposition source.

SECTION 13

Glossary

Term	Definition
ASR (Automatic Speech Recognition)	The process of converting audio speech into written text. Performed by a vendor API in Stage 1 of the pipeline.
Batch API	A vendor API mode that accepts thousands of requests in a single submission, processes them asynchronously overnight, and returns results in bulk at a substantially reduced cost versus synchronous calls.
Concurrency Controller	A software component that maintains exactly N concurrent in-flight API jobs at all times, preventing rate limit breaches while maximising throughput.
D0	The original call disposition submitted by the tele-caller after completing a call. This is the agent's own characterisation of the outcome.
D1	The AI-inferred call disposition derived from the actual transcript content by the LLM. D1 is compared to D0 to detect mismatches.
Dead-letter Queue	A holding queue for jobs that have failed all retry attempts. Items here are flagged for manual review and included in daily error reports.
Diarization	The process of separating a multi-speaker audio recording into individual speaker turns, labelling each segment with its speaker identity (AGENT / CUSTOMER).
Diarized Transcript	A text transcript where each utterance is attributed to a specific speaker, with start and end timestamps per turn.
Disposition	A classification of the outcome of a collection call — e.g., customer paid, promised to pay, refused, or was not reachable.
LLM (Large Language Model)	The AI model used to analyse transcripts and produce the four structured output modules. Accessed via a vendor API (Anthropic or OpenAI).
PII (Personally Identifiable Information)	Data that can identify a specific individual — names, phone numbers, account numbers. Not extracted or mapped separately by this system.
Script Adherence	The degree to which an agent followed the bank's prescribed calling stages in the correct sequence during a call.
Composite Score	The primary agent performance KPI — a weighted average of flow correctness, response handling, and call quality, scored 0–10 per call.